

# Practical

Start by loading usefull libraries:

```
library(dplyr)
library(INLA)
library(ggplot2)
library(patchwork)
library(inlabru)
# load some libraries to generate nice map plots
library(scico)
```

**Aim of this practical:** In this first practical we are going to look at some simple models

1. A Gaussian model with simulated data
2. A GLM model with random effects
3. A disease mapping model

we are going to learn:

- How to fit a simple model with `inlabru`
- How to explore the results
- How to change the prior distributions
- How to get predictions for missing data points

## Linear Model

As our first example we consider a simple linear regression model with Gaussian observations  $y_i \sim \mathcal{N}(\mu_i, \sigma^2)$ ,  $i = 1, \dots, N$  where  $\sigma^2$  is the observation error, and the mean parameter  $\mu_i$  is linked to the linear predictor through an identity function:

$$\eta_i = \mu_i = \beta_0 + \beta_1 x_i$$

where  $x_i$  is a covariate and  $\beta_0, \beta_1$  are parameters to be estimated.

To finalize the Bayesian model we need to assign a  $\text{Gamma}(a, b)$  prior to the precision parameter  $\tau = 1/\sigma^2$  and two independent Gaussian priors with mean 0 and precision  $\tau_\beta$  to the regression parameters  $\beta_0$  and  $\beta_1$ .

#### Question

What is the dimension of the hyperparameter vector and latent Gaussian field?

Answer

The hyperparameter vector has dimension 1,  $\boldsymbol{\theta} = (\tau)$  while the latent Gaussian field  $\mathbf{u} = (\beta_0, \beta_1)$  has dimension 2, 0 mean, and sparse precision matrix:

$$\mathbf{Q} = \tau_\beta \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$$

#### Note

We can write the linear predictor vector  $\boldsymbol{\eta} = (\eta_i, \dots, \eta_N)$  as

$$\boldsymbol{\eta} = \mathbf{A}\mathbf{u} = \mathbf{A}_1\mathbf{u}_1 + \mathbf{A}_2\mathbf{u}_2 = \begin{bmatrix} 1 \\ 1 \\ \vdots \\ 1 \end{bmatrix} \beta_0 + \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_N \end{bmatrix} \beta_1$$

Our linear predictor consists then of two components.

### Simulate example data

---

In this practical we will use simulated Gaussian data to get familiar with the `inlabru` workflow. Moreover, we will see how to change the prior distributions both for the fixed effects  $\beta_0$  and  $\beta_1$  and for the hyperparameter  $\tau = 1/\sigma^2$ . First, we simulate data from the model

$$y_i \sim \mathcal{N}(\eta_i, 0.1^2), \quad i = 1, \dots, 100$$

with

$$\eta_i = \beta_0 + \beta_1 x_i$$

where  $\beta_0 = 2, \beta_1 = 0.5$  and the values of the covariate  $x$  are generated from an `Uniform(0,1)` distribution. The simulated response and covariate data are then saved in a `data.frame` object.

```

beta = c(1,1)
sd_error = 1

n = 100
x = rnorm(n)
y = beta[1] + beta[2] * x + rnorm(n, sd = sd_error)

df = data.frame(y = y, x = x)

```

## Fitting a linear regression model with inlabru

---

### Defining model components

The model has two parameters to be estimated  $\beta_1$  and  $\beta_2$ . We need to define the two corresponding model components:

```

cmp = ~ Intercept(1) + beta_1(x, model = "linear")

```

The `cmp` object is here used to define model components. We can give them any useful names we like

#### Note

Note that `Intercept()` is one of `inlabru` special names and it is used to define a global intercept. You should explicitly exclude automatic intercept when not using the special `Intercept` name, e.g.

```

cmp = ~ -1 + myIntercept(1) + beta_1(x, model = "linear")

```

### Observation model construction

The next step is to construct the observation model by defining the model likelihood. The most important inputs here are the **formula**, the **family** and the **data**.

The **formula** defines how the components should be combined in order to define the model predictor.

```

formula = y ~ Intercept + beta_1

```

### **i** Note

In this case we can also use the shortcut `formula = y ~ .`. This will tell `inlarbu` that the model is linear and that it is not necessary to linearize the model and assess convergence.

The likelihood is defined using the `bru_obs()` function as follows:

```
lik = bru_obs(formula = y ~ .,
              family = "gaussian",
              data = df)
```

### **Fit the model**

We fit the model using the `bru()` functions which takes as input the components and the observation model:

```
fit.lm = bru(cmp, lik)
```

The `summary()` function will give access to some basic information about model fit and estimates

```
summary(fit.lm)
```

```
inlabru version: 2.12.0
```

```
INLA version: 25.02.10
```

```
Components:
```

```
Intercept: main = linear(1), group = exchangeable(1L), replicate = iid(1L), NULL
```

```
beta_1: main = linear(x), group = exchangeable(1L), replicate = iid(1L), NULL
```

```
Likelihoods:
```

```
Family: 'gaussian'
```

```
Tag: ''
```

```
Data class: 'data.frame'
```

```
Response class: 'numeric'
```

```
Predictor: y ~ .
```

```
Used components: effects[Intercept, beta_1], latent[]
```

```
Time used:
```

```
Pre = 0.376, Running = 0.182, Post = 0.0744, Total = 0.633
```

```
Fixed effects:
```

	mean	sd	0.025quant	0.5quant	0.975quant	mode	kld
Intercept	0.865	0.095	0.679	0.865	1.051	0.865	0
beta_1	0.967	0.088	0.795	0.967	1.139	0.967	0

Model hyperparameters:

	mean	sd	0.025quant	0.5quant
Precision for the Gaussian observations	1.14	0.161	0.848	1.13
			0.975quant	mode
Precision for the Gaussian observations		1.48	1.12	

Deviance Information Criterion (DIC) .....: 276.50

Deviance Information Criterion (DIC, saturated) ....: 105.46

Effective number of parameters .....: 3.00

Watanabe-Akaike information criterion (WAIC) ...: 276.56

Effective number of parameters .....: 2.95

Marginal log-Likelihood: -157.78

is computed


Posterior summaries for the linear predictor and the fitted values are computed

(Posterior marginals needs also 'control.compute=list(return.marginals.predictor=TRUE)')

We can see that both the intercept and slope and the error precision are correctly estimated.

We can then plot the marginal posterior for  $\beta_0$  as follows:

```
plot(fit.lm, "Intercept")
```



### Task

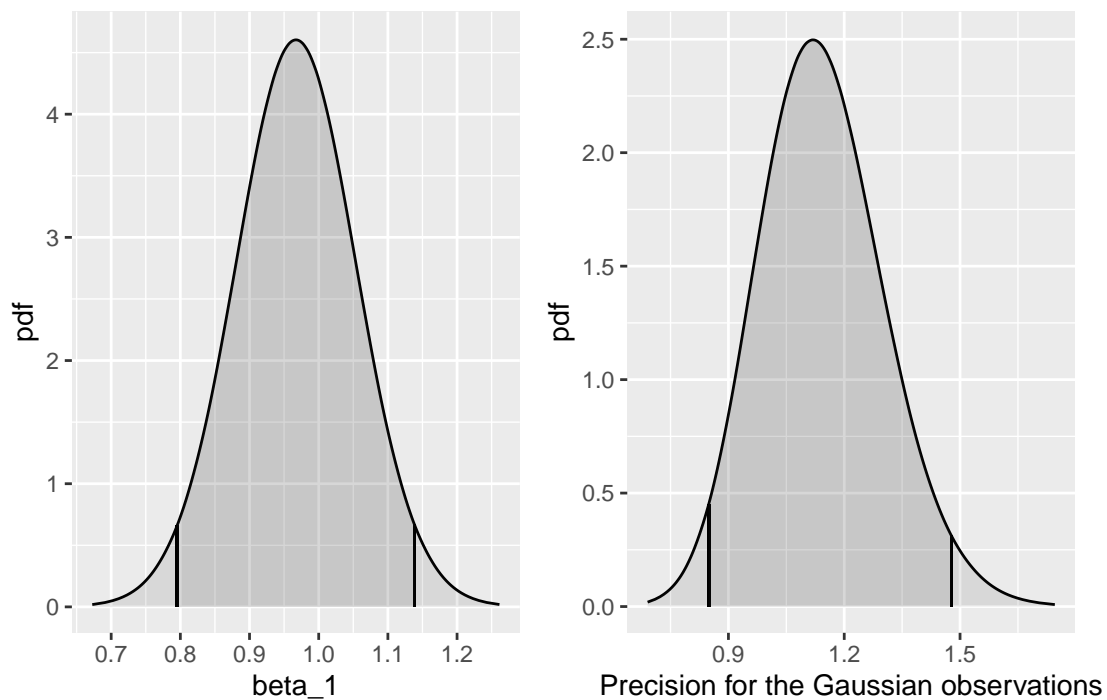
Plot the posterior marginals for  $\beta_1$  and for the precision of the observation error  $\pi(\tau|y)$

Take hint

See the `summary()` output to check the names for the different model components.

[Click here to see the solution](#)

```
plot(fit.lm, "beta_1") +  
plot(fit.lm, "Precision for the Gaussian observations")
```



### Task


Plot the fitted values with 95% Credible intervals.

Take hint

`bru` objects information about the linear predictor can be accessed through `fit.lm$summary.fitted.values`.

[Click here to see the solution](#)

```
df %>% mutate(post_mean = fit.lm$summary.fitted.values[1:100,"mean"],
              q25 = fit.lm$summary.fitted.values[1:100,"0.025quant"],
              q975 = fit.lm$summary.fitted.values[1:100,"0.975quant"]))%>%
  ggplot()+geom_point(aes(x=x,y=y),alpha=0.5,color="grey40")+
  geom_line(aes(x=x,y=post_mean),col=2)+
  geom_ribbon(aes(x = x, ymax = q975, ymin = q25),fill="tomato", alpha = 0.3)
```



## Generate model predictions

Now we can take the fitted `bru` object and use the `predict` function to produce predictions given a new set of values for the model covariates or the original values used for the model fit

```
new_data = data.frame(x = c(df$x, runif(10)),
                      y = c(df$y, rep(NA,10)))
pred = predict(fit.lm, new_data, ~ Intercept + beta_1)
```

## Plot




Figure 1: Data and 95% credible intervals

## R Code

```
pred %>% ggplot() +  
  geom_point(aes(x,y), alpha = 0.3) +  
  geom_line(aes(x,mean)) +  
  geom_line(aes(x, q0.025), linetype = "dashed")+  
  geom_line(aes(x, q0.975), linetype = "dashed")+  
  xlab("Covariate") + ylab("Observations")
```